

## WEEK – 4

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

dataset = pd.read_csv('Salary_Data.csv')
dataset.head()

# data preprocessing

X = dataset.iloc[:, :-1].values #independent variable array
y = dataset.iloc[:, 1].values #dependent variable vector

# splitting the dataset

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=1/3,random_state=0)

# fitting the regression model

from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train,y_train) #actually produces the linear eqn for the data

# predicting the test set results
y_pred = regressor.predict(X_test)

y_pred
y_test

# visualizing the results

#plot for the TRAIN

plt.scatter(X_train, y_train, color='red') # plotting the observation line
plt.plot(X_train, regressor.predict(X_train), color='blue') # plotting the regression line
plt.title("Salary vs Experience (Training set)") # stating the title of the graph
plt.xlabel("Years of experience") # adding the name of x-axis
plt.ylabel("Salaries") # adding the name of y-axis
plt.show() # specifies end of graph
```

```
#plot for the TEST
```

```
plt.scatter(X_test, y_test, color='red')
```

```
plt.plot(X_train, regressor.predict(X_train), color='blue') # plotting the regression line
```

```
plt.title("Salary vs Experience (Testing set)")
```

```
plt.xlabel("Years of experience")
```

```
plt.ylabel("Salaries")
```

```
plt.show()
```

WEEK – 5

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
from sklearn import preprocessing
```

```
# importing data
```

```
df = pd.read_csv('Real-estate1.csv')
```

```
df.drop('No', inplace=True, axis=1)
```

```
print(df.head())
```

```
print(df.columns)
```

```
# plotting a scatterplot
```

```
sns.scatterplot(x='X4 number of convenience stores',
```

```
y='Y house price of unit area', data=df)

# creating feature variables
X = df.drop('Y house price of unit area', axis=1)
y = df['Y house price of unit area']
print(X)
print(y)

# creating train and test sets
X_train, X_test, y_train, y_test = train_test_split(
X, y, test_size=0.3, random_state=101)

# creating a regression model
model = LinearRegression()

# fitting the model
model.fit(X_train, y_train)

# making predictions
predictions = model.predict(X_test)

# model evaluation
print('mean_squared_error : ', mean_squared_error(y_test, predictions))

print('mean_absolute_error : ', mean_absolute_error(y_test, predictions))
```

## WEEK – 6

```
import numpy as np
import pandas as pd
```

```

from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

import matplotlib.pyplot as plt

from sklearn.tree import DecisionTreeClassifier, plot_tree

def importdata():

    balance_data = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-' +
'databases/balance-scale/balance-scale.data',sep=',', header=None)

    print("Dataset Length: ", len(balance_data))

    print("Dataset Shape: ", balance_data.shape)

    print("Dataset: ", balance_data.head())

    return balance_data


def splitdataset(balance_data):

# Separating the target variable

    X = balance_data.values[:, 1:5]

    Y = balance_data.values[:, 0]

    X_train, X_test, y_train, y_test = train_test_split(

        X, Y, test_size=0.3, random_state=100)

    return X, Y, X_train, X_test, y_train, y_test

def train_using_gini(X_train, X_test, y_train):

# Creating the classifier object

    clf_gini = DecisionTreeClassifier(criterion="gini",random_state=100,max_depth=3,
min_samples_leaf=5)

    clf_gini.fit(X_train, y_train)

    return clf_gini

def train_using_entropy(X_train, X_test, y_train):

# Decision tree with entropy

```

```

    clf_entropy = DecisionTreeClassifier(criterion="entropy",
random_state=100,max_depth=3, min_samples_leaf=5)

    clf_entropy.fit(X_train, y_train)

    return clf_entropy

# Function to make predictions
def prediction(X_test, clf_object):

    y_pred = clf_object.predict(X_test)

    print("Predicted values:")

    print(y_pred)

    return y_pred

# Placeholder function for cal_accuracy
def cal_accuracy(y_test, y_pred):

    print("Confusion Matrix: ",confusion_matrix(y_test, y_pred))

    print("Accuracy : ",accuracy_score(y_test, y_pred)*100)

    print("Report : ",classification_report(y_test, y_pred))

# Function to plot the decision tree
def plot_decision_tree(clf_object, feature_names, class_names):

    plt.figure(figsize=(15, 10))

    plot_tree(clf_object, filled=True,
feature_names=feature_names,class_names=class_names, rounded=True)

    plt.show()

if __name__ == "__main__":

    data = importdata()

    X, Y, X_train, X_test, y_train, y_test = splitdataset(data)

    clf_gini = train_using_gini(X_train, X_test, y_train)

    clf_entropy = train_using_entropy(X_train, X_test, y_train)

    plot_decision_tree(clf_gini, ['X1', 'X2', 'X3', 'X4'], ['L', 'B', 'R'])

    plot_decision_tree(clf_entropy, ['X1', 'X2', 'X3', 'X4'], ['L', 'B', 'R'])

    print("Results Using Gini Index:")

```

```
y_pred_gini = prediction(X_test, clf_gini)
cal_accuracy(y_test, y_pred_gini)
```

## WEEK – 7

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
```

```
names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'Class']
```

```
df = pd.read_csv(url, names=names)
df
```

```
X = df.iloc[:, :-1].values
```

```
y = df.iloc[:, 4].values
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
scaler.fit(X_train)
```

```
X_train = scaler.transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
```

```
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

## WEEK – 8

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# load dataset
diab_df = pd.read_csv("diabetes.csv")
diab_df.head()

#split dataset in features and target variable
diab_cols = ['Pregnancies', 'Insulin', 'BMI',
'Age','Glucose','BloodPressure','DiabetesPedigreeFunction']
X = diab_df[diab_cols]# Features
y = diab_df.Outcome # Target variable

# Splitting Data
```

```

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=0)

# Model Development and Prediction and instantiate the model
logreg = LogisticRegression(solver='liblinear')

# fit the model with data
logreg.fit(X_train,y_train)

# predicting
y_pred=logreg.predict(X_test)

y_pred

# Model Evaluation using Confusion Matrix
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)

cnf_matrix

# Visualizing Confusion Matrix using Heatmap
class_names=[0,1] # name of classes

fig, ax = plt.subplots()

tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)

# create heatmap
sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')

# Confusion Matrix Evaluation Metrics
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))

```



```
print("Precision:",metrics.precision_score(y_test, y_pred))  
print("Recall:",metrics.recall_score(y_test, y_pred))
```

## WEEK – 9

```
import numpy as np  
from sklearn.datasets import make_blobs  
class KMeans:  
    def _init_(self, n_clusters, max_iters=100):  
        self.n_clusters = n_clusters  
        self.max_iters = max_iters  
    def fit(self, X):  
        self.centroids = X[np.random.choice(X.shape[0], self.n_clusters, replace=False)]  
        for _ in range(self.max_iters):  
            labels = self._assign_labels(X)  
            new_centroids = self._update_centroids(X, labels)  
            if np.all(self.centroids == new_centroids):  
                break  
            self.centroids = new_centroids  
    def _assign_labels(self, X):  
        distances = np.linalg.norm(X[:, np.newaxis] - self.centroids, axis=2)  
        return np.argmin(distances, axis=1)  
    def _update_centroids(self, X, labels):  
        new_centroids = np.array([X[labels == i].mean(axis=0)  
                                   for i in range(self.n_clusters)])  
        return new_centroids  
X, _ = make_blobs(n_samples=300, centers=3, random_state=42)  
kmeans = KMeans(n_clusters=3)
```

```
kmeans.fit(X)

labels = kmeans._assign_labels(X)

print("Cluster Assignments:", labels)

print("Final Centroids:", kmeans.centroids)
```

## EXPERIMENT 10 ML

```
import pandas as pd

from pandas.plotting import scatter_matrix

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split, KFold, cross_val_score

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

from sklearn.linear_model import LogisticRegression

from sklearn.tree import DecisionTreeClassifier

from sklearn.neighbors import KNeighborsClassifier

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

from sklearn.naive_bayes import GaussianNB

from sklearn.svm import SVC
```

# 1. Load dataset

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
```

```
names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
```

```
dataset = pd.read_csv(url, names=names)
```

# 2. Summarize the Dataset

# Shape

```
print(dataset.shape)
```

# Plotting a scatterplot

```
sns.scatterplot(x='sepal-length', y='petal-length', data=dataset)
```

```
plt.show()
```

# Creating feature variables

```
X = dataset.drop('class', axis=1)
```

```
y = dataset['class']
```

```
print(X)
```

```
print(y)
```

# Creating train and test sets

```
validation_size = 0.20
```

```
seed = 7
```

```
X_train, X_validation, Y_train, Y_validation = train_test_split(X, y,  
test_size=validation_size, random_state=seed)
```

# Test Harness

```
seed = 7
```

```
scoring = 'accuracy'
```

```

# Build and evaluate models

models = []

models.append(('LR', LogisticRegression()))

models.append(('LDA', LinearDiscriminantAnalysis()))

models.append(('KNN', KNeighborsClassifier()))

models.append(('CART', DecisionTreeClassifier()))

models.append(('NB', GaussianNB()))

models.append(('SVM', SVC()))


# Evaluate each model in turn

results = []

names = []

for name, model in models:

    kfold = KFold(n_splits=10, random_state=seed, shuffle=True)

    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold, scoring=scoring)

    results.append(cv_results)

    names.append(name)

    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())

    print(msg)


# Compare Algorithms

fig = plt.figure()

fig.suptitle('Algorithm Comparison')

ax = fig.add_subplot(111)

plt.boxplot(results)

ax.set_xticklabels(names)

plt.show()

```

```
# Multivariate Plots
```

```
# Scatter plot matrix
```

```
scatter_matrix(dataset)
```

```
plt.show()
```

```
# Box plot
```

```
dataset.plot(kind='box', subplots=True, layout=(2, 2), sharex=False, sharey=False)
```

```
plt.show()
```

```
# Histogram
```

```
dataset.hist()
```

```
plt.show()
```

```
# Descriptive statistics
```

```
print(dataset.describe())
```

```
# Class distribution
```

```
print(dataset.groupby('class').size())
```